

Introduction to REST APIs

Olivier Liechti & Laurent Prévost
COMEM Web Services

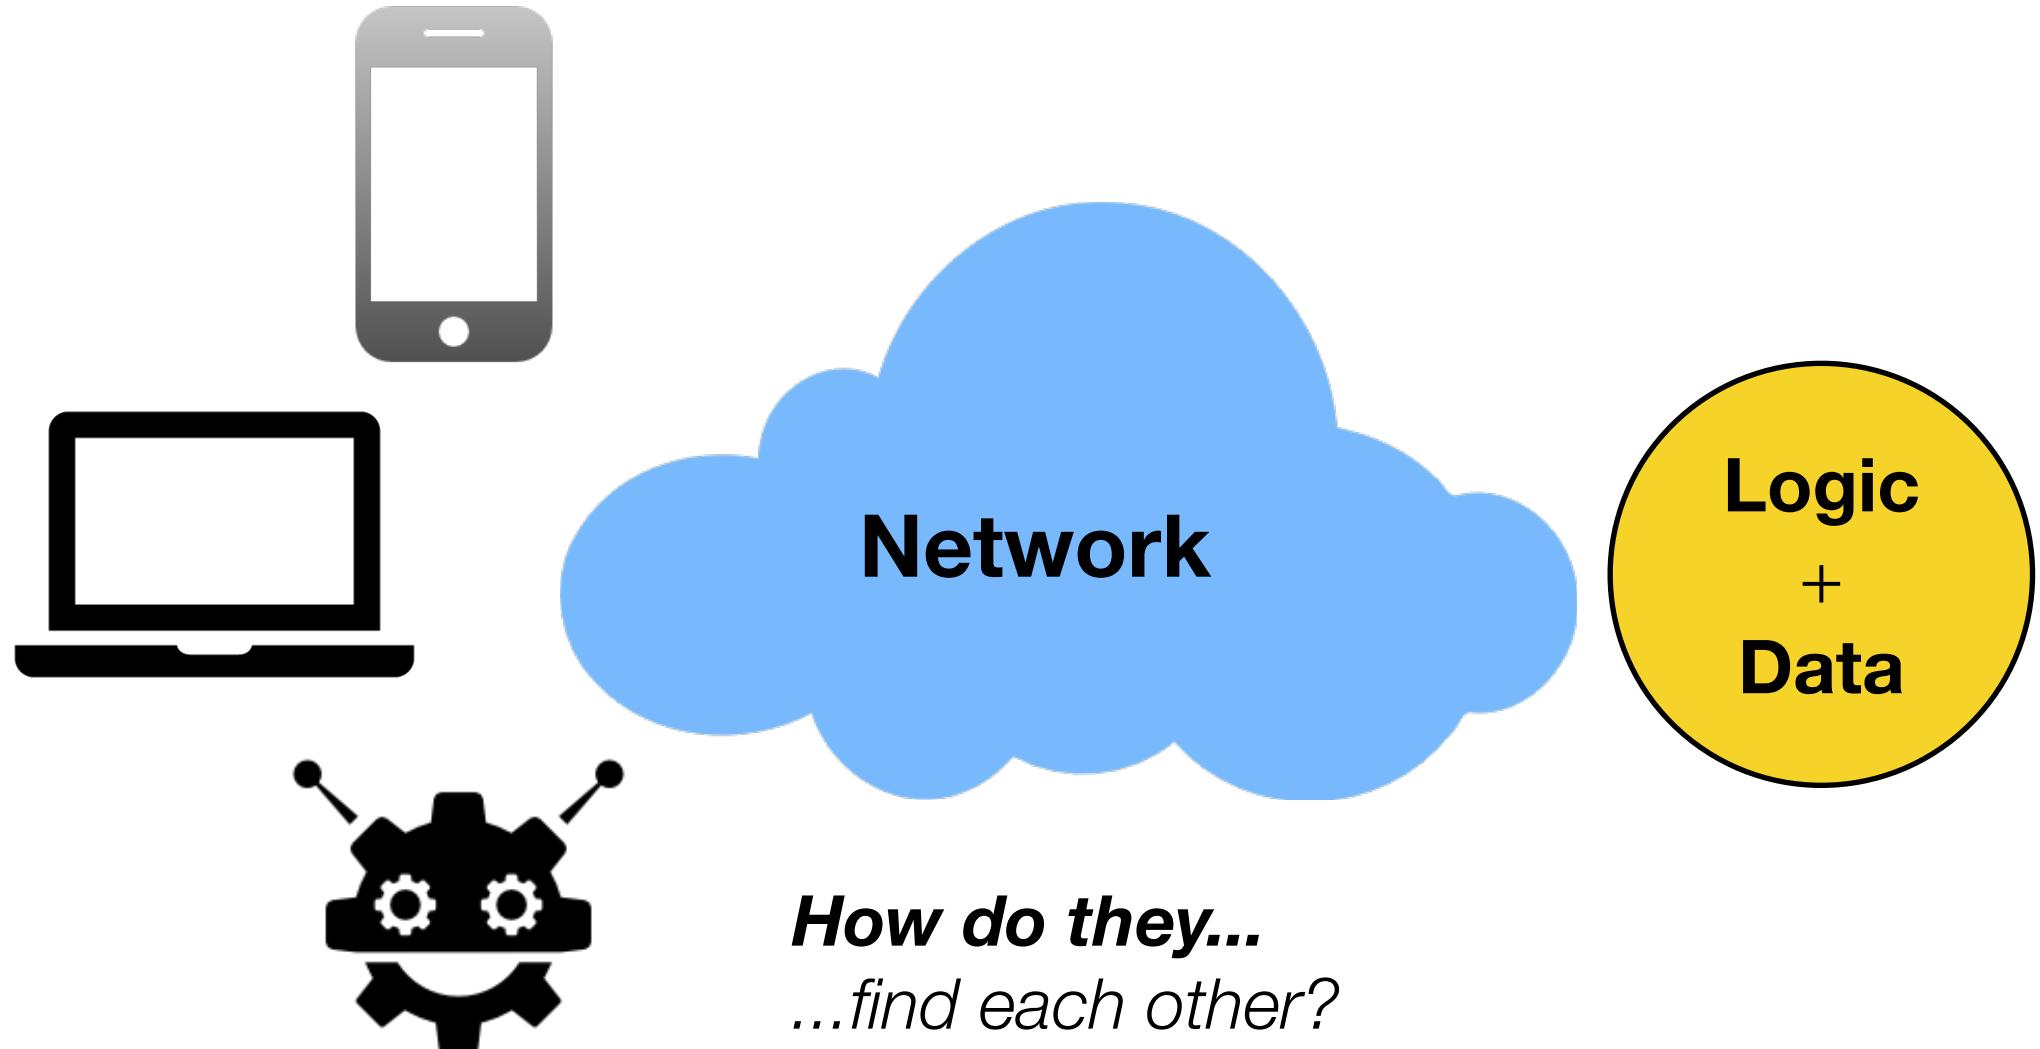


Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud



Big Web Services vs REST

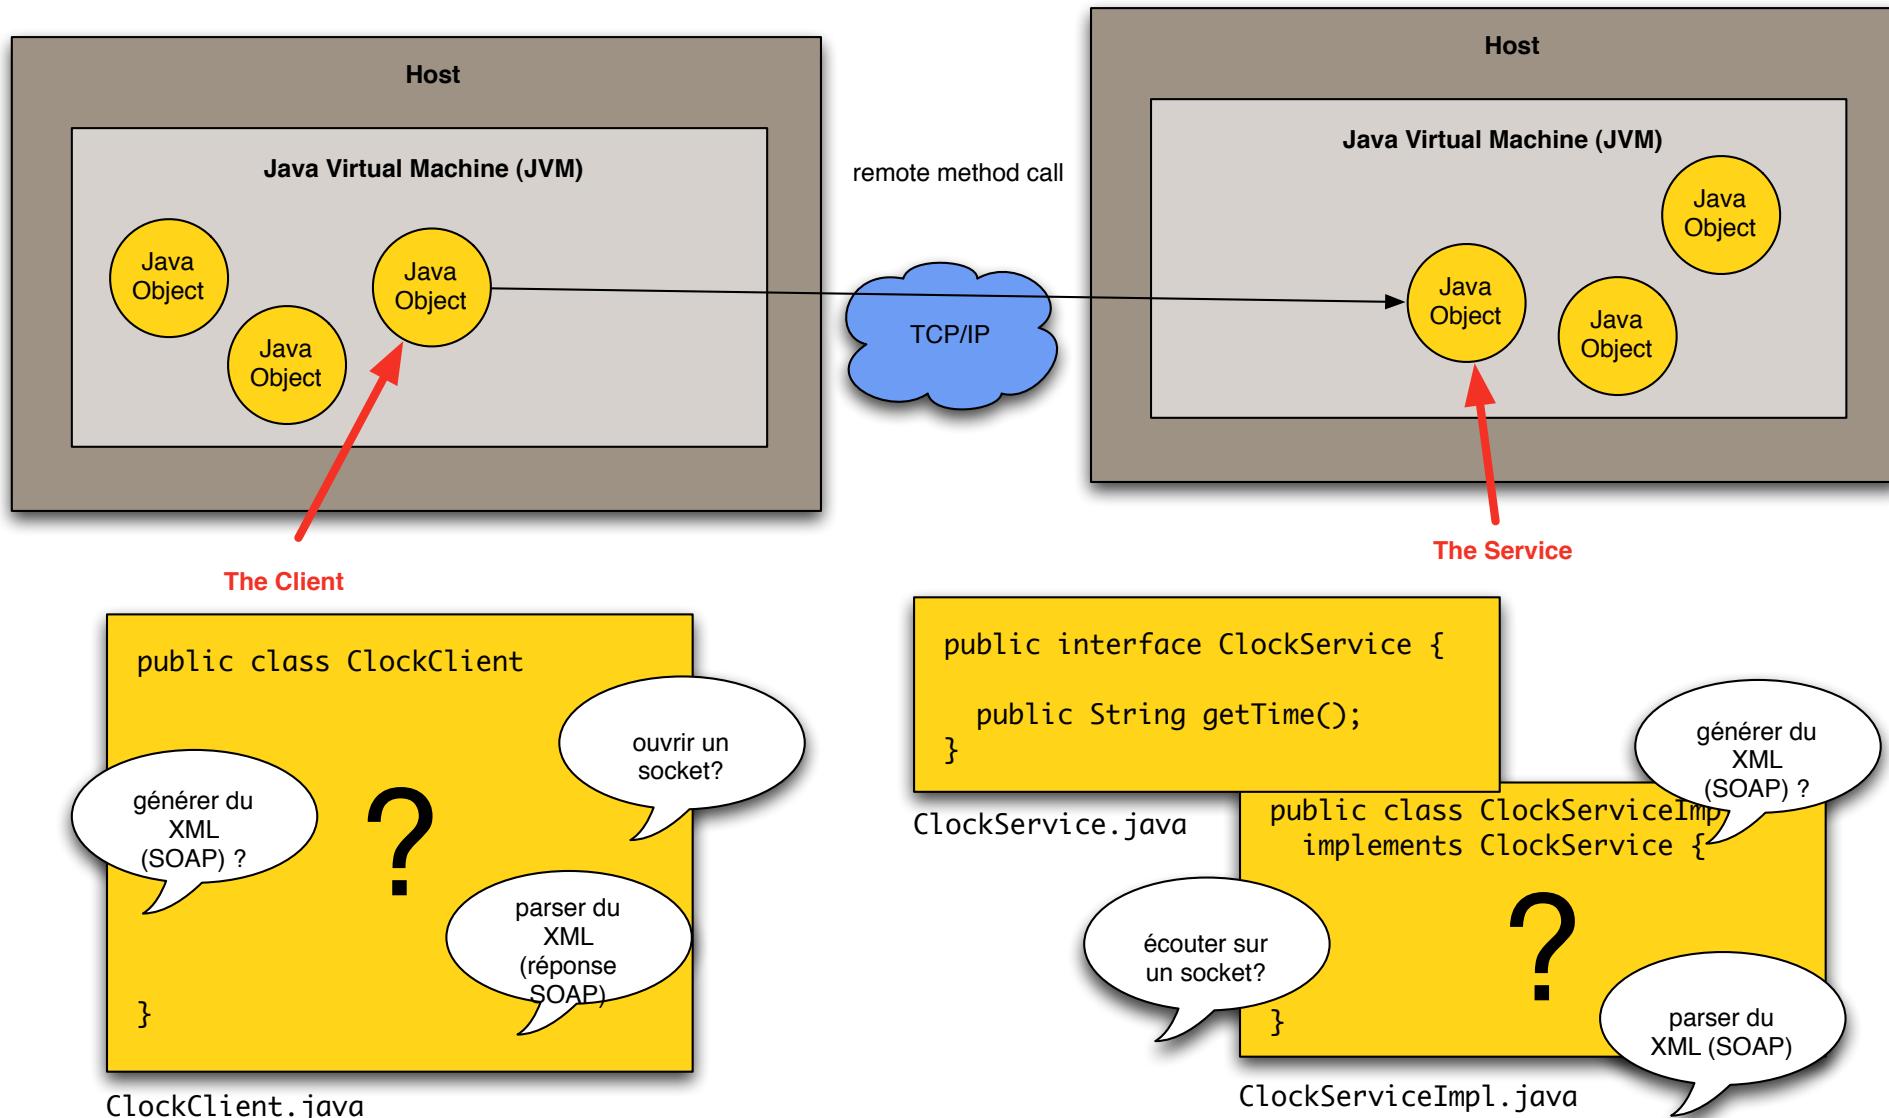
What is a Web Service?



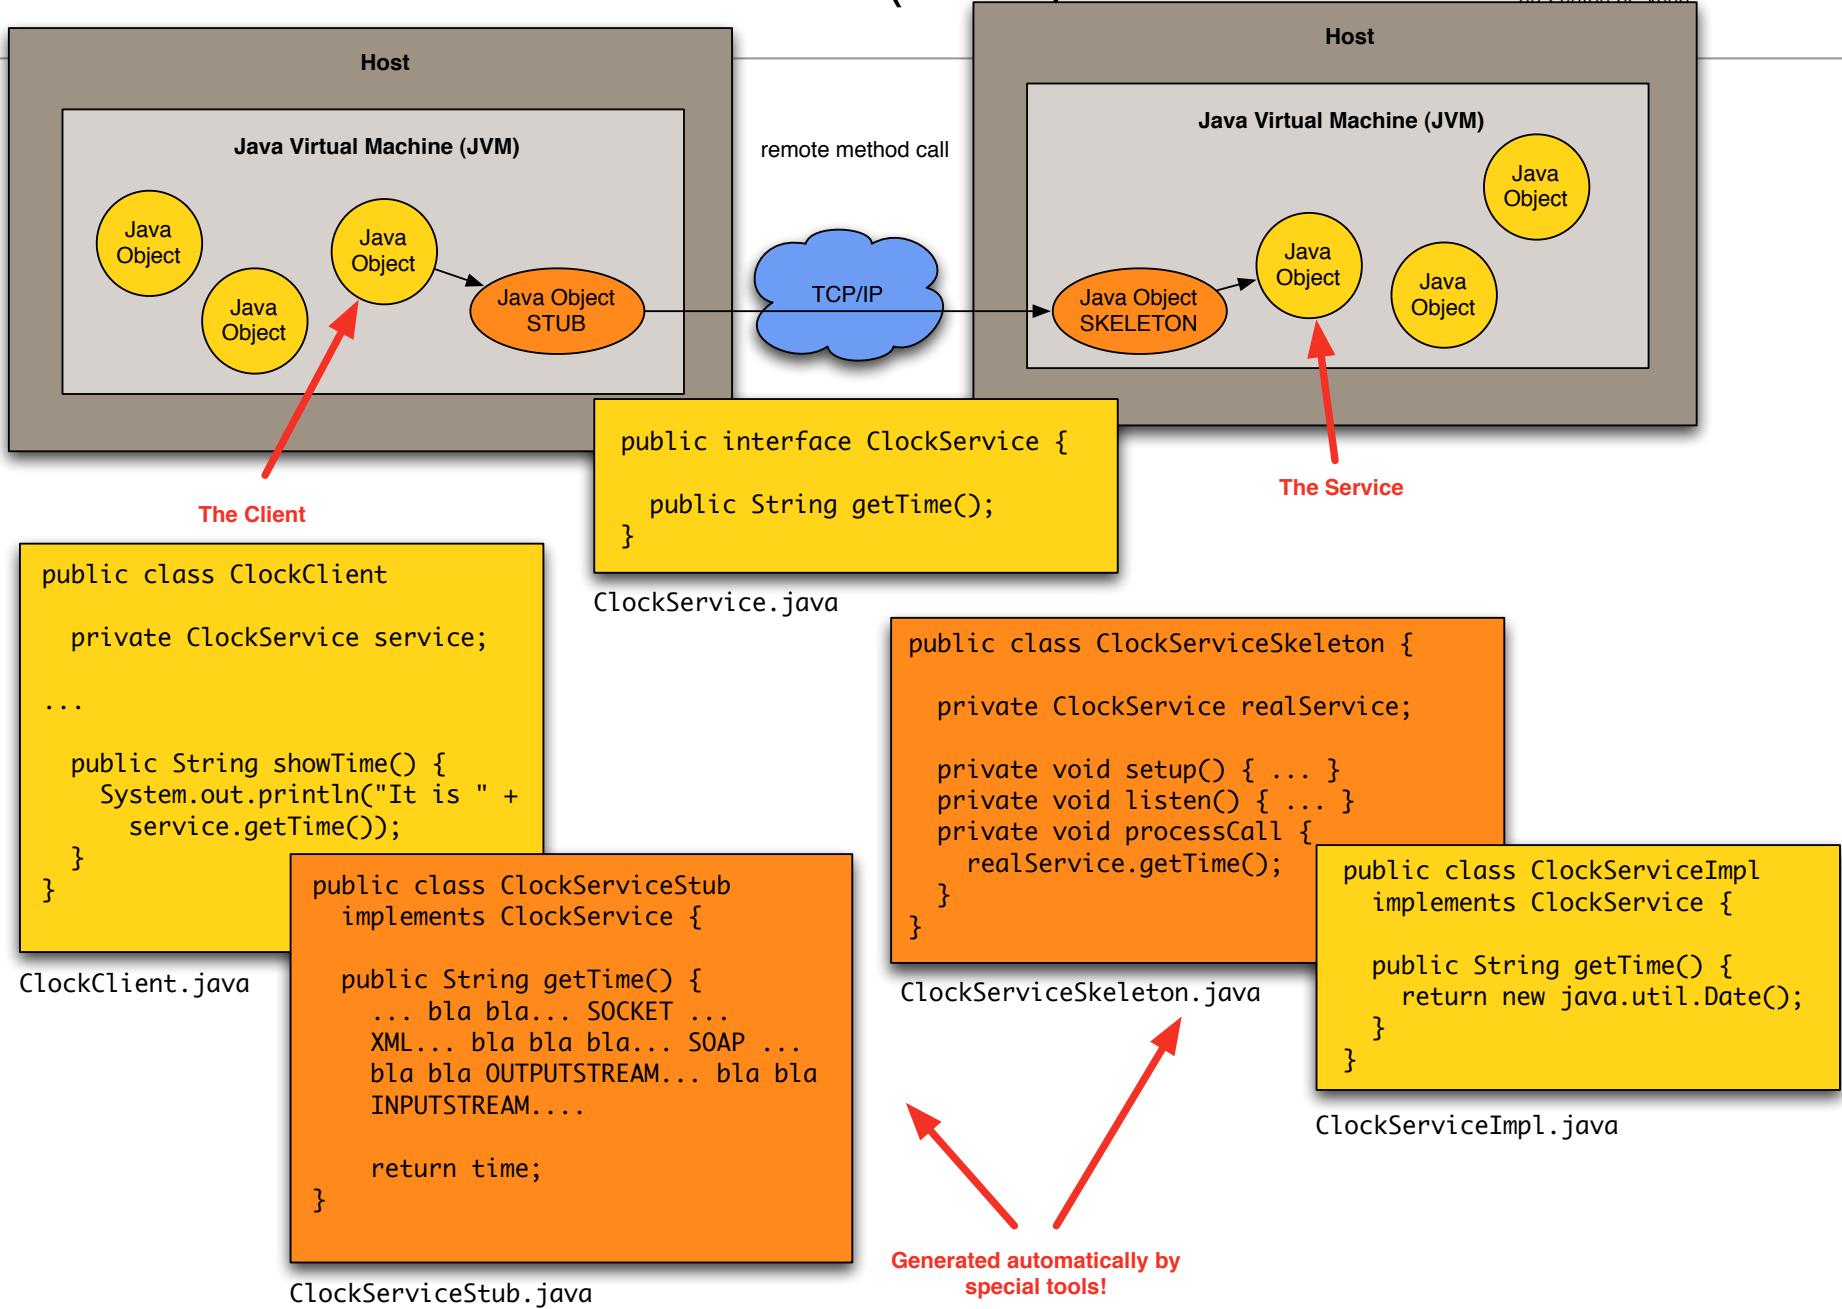
How do they...

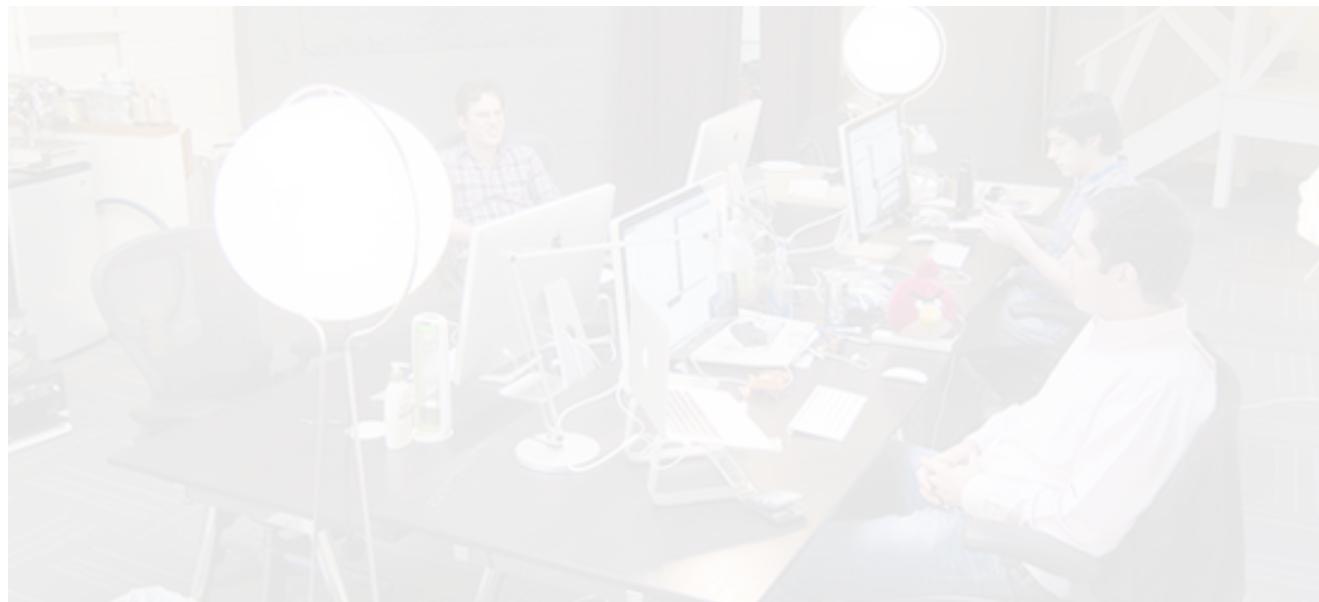
- ...find each other?*
- ...know what logic can be invoked?*
- ...talk to each other?*

Remote services in Java (RMI)



Remote services in Java (RMI)



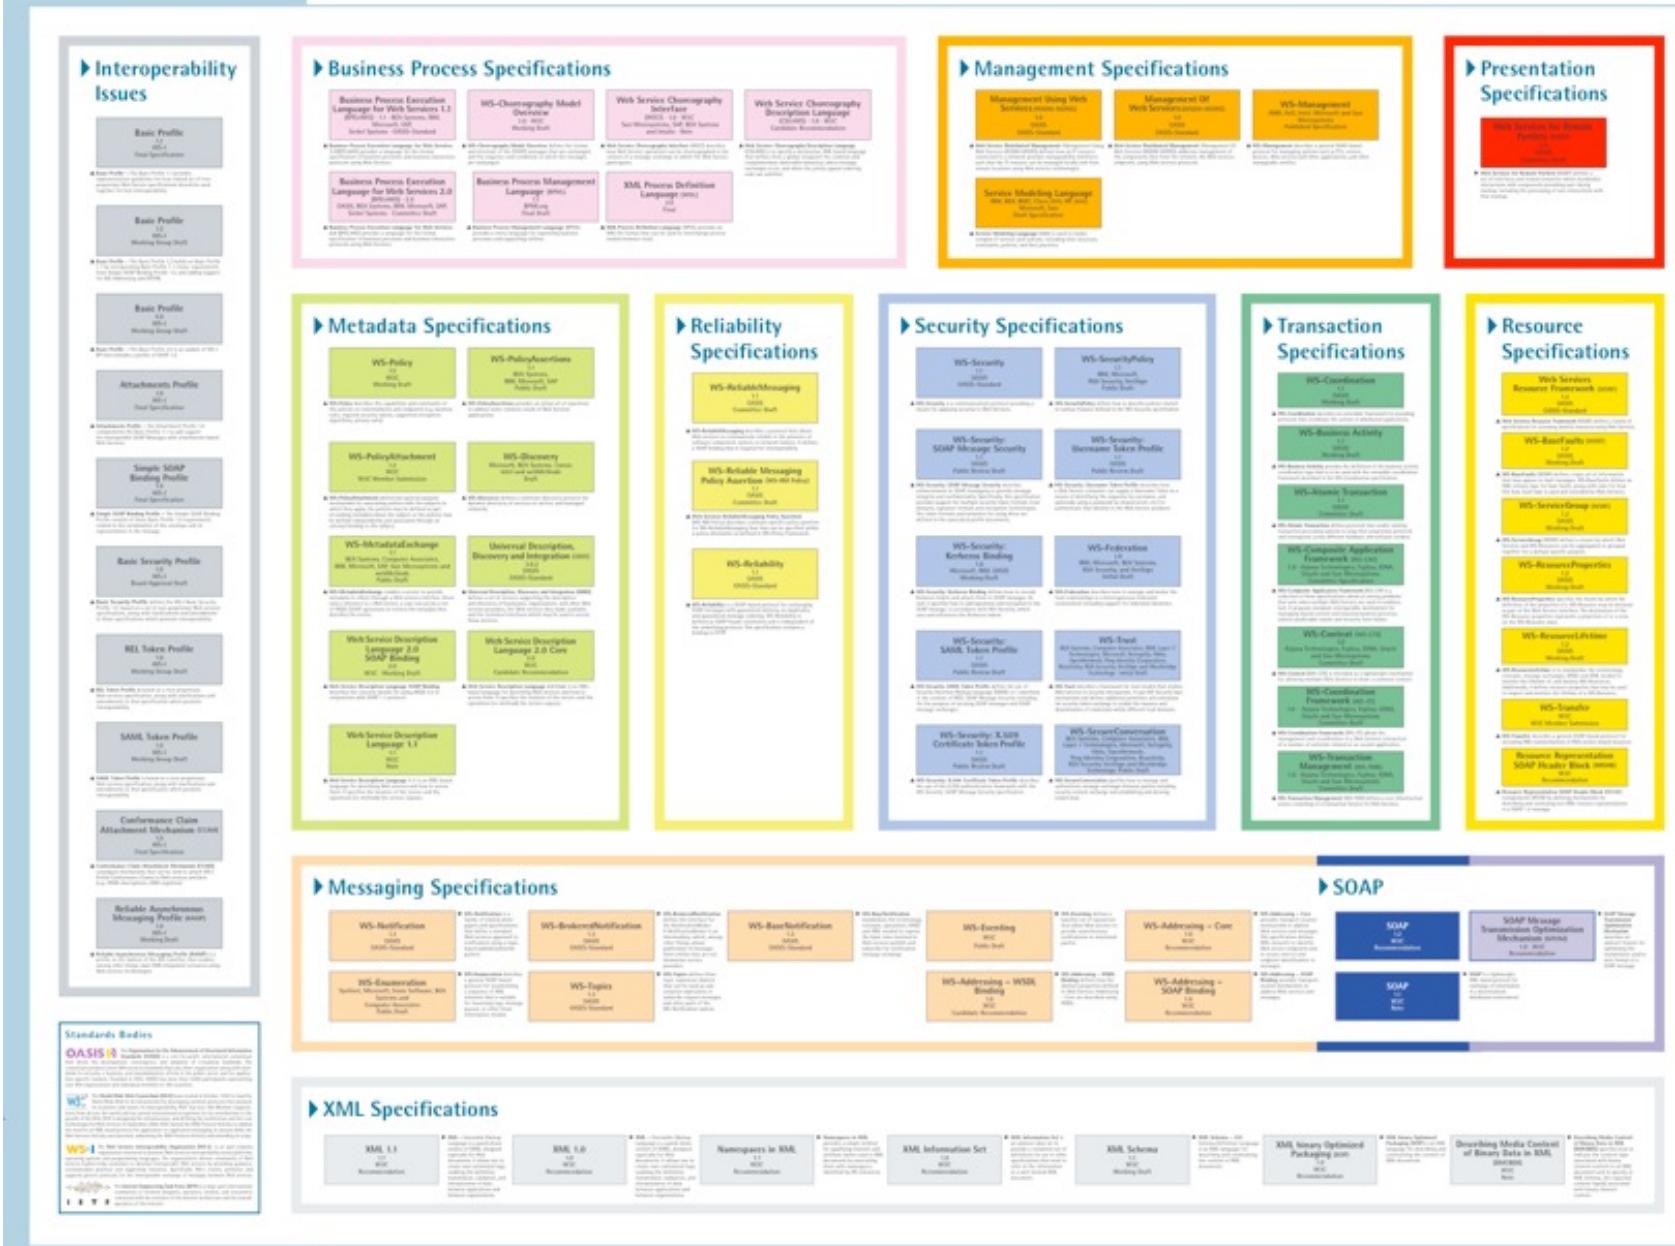


The “Big” Web Services Approach

Big Web Services

- **Approach**
 - Services are often designed and developed with a **RPC style** (even if Document-Oriented Services are possible).
- **Core Standards**
 - Simple Object Access Protocol (**SOAP**)
 - Web Services Description Language (**WSDL**)
- **Benefits**
 - **Very rich protocol stack** (support for security, transactions, reliable transfer, etc.)
- **Problem**
 - **Very rich protocol stack** (complexity, verbosity, incompatibility issues, theoretical human readability, etc.)

Web Services Standards Overview



innoQ

innoQ Deutschland GmbH
Hahnstraße 17
D-40880 Ratingen
Phone +49 2102 77 162 - 100
info@innoq.com • www.innoq.com

www.Q-Schweiz.ch
Gewerbestrasse 11
CH-6330 Cham
Phone +41 41 243 06 11

Web Services Standards Overview



innoQ

immobilienQ Deutschland GmbH
Hahnstraße 17
D-40880 Ratingen
Phone +49 2182 37 162-100
info@immobilienQ.com · www.immobilienQ.com

maul® Schweiz GmbH
Gewerbestrasse 11
CH-6330 Cham
Phone +41 41 243 06 00



Enterprise Platforms (Java EE, .NET, etc.) hide some of the complexity.

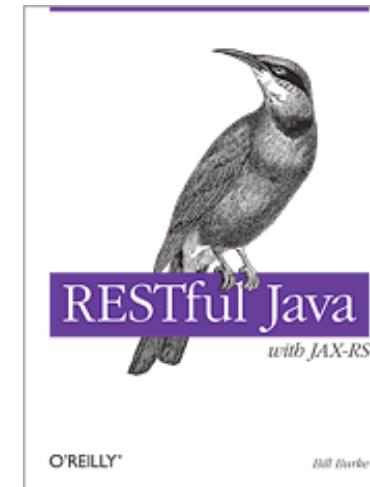
```
@Stateless  
@WebService  
public class Demo implements DemoLocal {  
  
    @Override  
    public String getTime() {  
        return new Date().toString();  
    }  
  
    @Override  
    public long computeSum(long v1, long v2) {  
        return v1 + v2;  
    }  
}
```

Adding a **@WebService** annotation does the magic. The application server takes care of all the gory details: generation of the WSDL interface, marshalling/unmarshalling of the SOAP messages. Still...



The REST Approach

RESTful Web Services



The REST Architectural Style

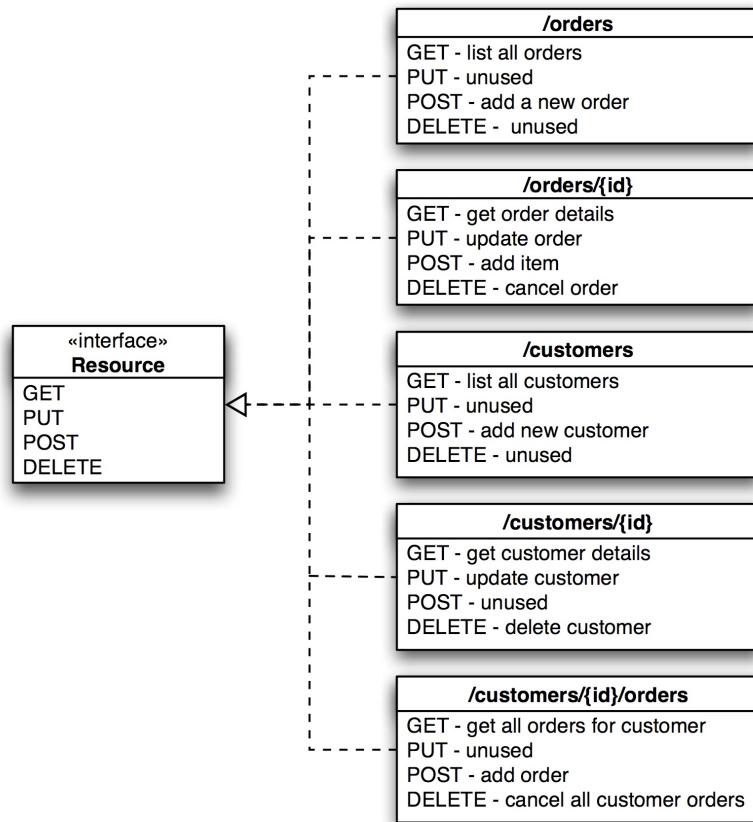
- REST: **REpresentational State Transfer**
- REST is an **architectural style** for building distributed systems.
- REST has been introduced in **Roy Fielding's Ph.D. thesis** (Roy Fielding has been a contributor to the HTTP specification, to the apache server, to the apache community).
- The WWW is **one example** for a distributed system that exhibits the characteristics of a REST architecture.

Principles of a REST Architecture

- The state of the application is captured in a **set of resources**
 - Users, photos, comments, tags, albums, etc.
- Every resource is **identified with a standard format** (e.g. URL)
- Every resource can have **several representations**
- There is one **unique interface for interacting** with resources (e.g. HTTP methods)

References

- Very good article, with presentation of key concepts and illustrative examples:
 - <http://www.infoq.com/articles/rest-introduction>
- Suggestions for the design of “pragmatic APIs”
 - <http://www.vinaysahni.com/best-practices-for-a-pragmatic-restful-api>



HTTP is a protocol for interacting with "**resources**"

What is a “Resource”

- At first glance, one could think that a “resource” is a file on a web server:
 - an HTML document, an XML document, a PNG document
- That fits the vision of the “static content” web
- But of course, the web is now more than a huge library of hypermedia documents:
 - through the web, we interact with services and a lot of the content is dynamic.
 - more and more, through the web we interact with physical objects (machines, sensors, actuators)
 - We need a more generic definition for resources!

What is a “Resource”?

- A resource is "something" that can be named and uniquely identified:
 - Example 1: an article published in the "24 heures" newspaper
 - Example 2: the collection of articles published in the sport section of the newspaper
 - Example 3: a person's resume
 - Example 4: the current price of the Nestlé stock quote
 - Example 5: the vending machine in the school hallway
 - Example 6: the list of grades of the student Jean Dupont
- URL (Uniform Resource Locator) is a mechanism for identifying resources
 - Exemple 1: <http://www.24heures.ch/vaud/vaud/2008/08/04/trente-etudiants-partent-rencontre-patrons>
 - Exemple 2: <http://www.24heures.ch/articles/sport>
 - Exemple 5: <http://www.smart-machines.ch/customers/heig/machines/8272>

Resource vs. Representation

- A "resource" can be something intangible (stock quote) or tangible (vending machine)
- The HTTP protocol supports the exchange of data between a client and a server.
- Hence, what is exchanged between a client and a server is **not** the resource. It is a **representation** of a resource.
- Different representations of the same resource can be generated:
 - HTML representation
 - XML representation
 - PNG representation
 - WAV representation
- **HTTP provides the content negotiation mechanisms!!**

CRUD

Create

Read

Update

Delete

CREATE RUD

heig-vd

Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud

```
POST /users/ HTTP/1.1
Host: www.myservice.com
Content-type: application/json
```

```
{
  "firstName" : "olivier",
  "lastName" : "liechti"
}
```

```
HTTP/1.1 201 Created
Location: /users/7823
```

```
GET /users/ HTTP/1.1
Host: www.myservice.com
Accept: application/json
```

```
HTTP/1.1 200 OK
Content-type: application/json
```

```
[
  { "id" : 7823 },
  {"id" : 7824}
]
```

```
GET /users/7823 HTTP/1.1
Host: www.myservice.com
Accept: application/json
```

```
HTTP/1.1 200 OK
Content-type: application/json
```

```
{
  "firstName": "olivier",
  "lastName": "liechti"
}
```

CR UPDATE D

```
PUT /users/7823 HTTP/1.1
Host: www.myservice.com
Content-type: application/json
```

```
{
  "firstName" : "olivier",
  "lastName" : "Liechti"
}
```

```
HTTP/1.1 204 No Content
```

```
DELETE /users/7823 HTTP/1.1
Host: www.myservice.com
```

```
HTTP/1.1 204 No Content
```

References in resource representations

```
GET /users/7823 HTTP/1.1
Host: www.myservice.com
```

```
HTTP/1.1 200 OK
Content-type: application/json

{
  'firstName' : 'olivier',
  'lastName' : 'liechti',
  'organizationId' : 892
}
```

ID

```
HTTP/1.1 200 OK
Content-type: application/json

{
  'firstName' : 'olivier',
  'lastName' : 'liechti',
  'organization' : '/organizations/892'
}
```

URL

Object graphs & URLs

```
GET /companies/8939/employees/192/office HTTP/1.1
```

VS

```
GET /office?occupantId=192 HTTP/1.1
```

```
GET /cars/882/wheels/frontLeft HTTP/1.1
```

VS

```
GET /wheels/89278927 HTTP/1.1
```

Filtering & Sorting in collections

```
GET /users/?filterByZipCode=1446&orderBy=lastName HTTP/1.1
Host: www.myservice.com
```

Pagination in collections

```
GET /users/ HTTP/1.1
Host: www.myservice.com
Accept: application/json
X-Page-Size: 10
```

```
HTTP/1.1 200 OK
Content-type: application/json
X-Page-Number: 1
X-Elements-Count: 23
X-Pages-Count: 3
X-Page-Size: 10

[{}, ...]
```

```
GET /users/ HTTP/1.1
Host: www.myservice.com
Accept: application/json
X-Page-Size: 10
X-Page-Number: 2
```

Alternative: link headers

<http://tools.ietf.org/html/rfc5988#page-6>



<https://developer.github.com/guides/traversing-with-pagination/>

Link: <<https://api.github.com/search/code?q=addClass+user%3Amozilla&page=15>>; rel="next",
<<https://api.github.com/search/code?q=addClass+user%3Amozilla&page=34>>; rel="last",
<<https://api.github.com/search/code?q=addClass+user%3Amozilla&page=1>>; rel="first",
<<https://api.github.com/search/code?q=addClass+user%3Amozilla&page=13>>; rel="prev"

URLs, actions & verbs

GET /cars/9383/startEngine HTTP/1.1



VS

POST /cars/9383/startEngine HTTP/1.1

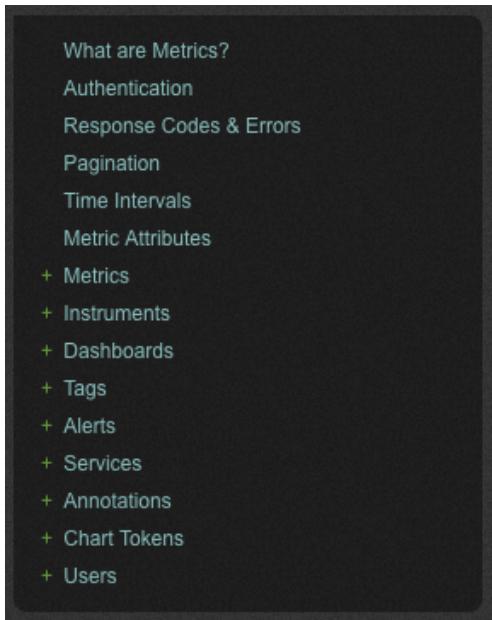
VS

POST /cars/9383/commands HTTP/1.1

Content-type: application/json

```
{  
  'command' : 'startEngine',  
  'params' : [  
    'when' : 'ASAP',  
    'how' : 'withTurbo'  
  ]  
}
```

Look at Some Examples



Documentation

Getting Started

API Reference

- Overview
- Authentication
- Things
- Properties
- Locations
- Products
- Collections
- Redirection Service
- Search

Code Examples

Documentation

Search Documentation

Overview

Authentication

Real-time

iPhone Hooks

API Console

Endpoints

- Users
- Relationships
- Media
- Comments
- Likes
- Tags
- Locations
- Geographies

Embedding

Libraries

Forum

<http://dev.librato.com/v1>

[https://dev.evrythng.com/
documentation/api](https://dev.evrythng.com/documentation/api)



[http://instagram.com/
developer/endpoints/](http://instagram.com/developer/endpoints/)



Instagram

Short description of the resource (domain model)

What Are Metrics?

Metrics are custom measurements stored in Librato's Metrics service. These measurements are created and may be accessed programmatically through a set of RESTful API calls. There are currently two types of metrics that may be stored in Librato Metrics, **gauges** and **counters**.

Gauges

Gauges capture a series of measurements where each measurement represents the value under observation at one point in time. The value of a gauge typically varies between some known minimum and maximum. Examples of gauge measurements include the requests/second serviced by an application, the amount of available disk space, the current value of \$AAPL, etc.

Counters

Counters track an increasing number of occurrences of some event. A counter is unbounded and always monotonically increasing in any given run. A new run is started anytime that counter is reset to zero. Examples of counter measurements include the number of connections made to an app, the number of visitors to a website, the number of times a write operation failed, etc.

Metric Properties

Some common properties are supported across all types of metrics:

`name`

Each metric has a name that is unique to its class of metrics e.g. a gauge name must be unique. The name identifies a metric in subsequent API calls to store/query individual measurements. The name can be up to 63 characters in length. Valid characters for metric names are 'A-Za-z0-9:_-'.

`period`

The `period` of a metric is an integer value that describes (in seconds) the standard reporting interval for the metric. Setting the period enables Metrics to detect abnormal interruptions in reporting and automatically resume reporting.

navigator

What are Metrics?

Authentication

Response Codes & Errors

Pagination

Time Intervals

Metric Attributes

- Metrics

`GET /metrics`

`POST /metrics`

`DELETE /metrics`

`GET /metrics/:name`

`PUT /metrics/:name`

`DELETE /metrics/:name`

+ Instruments

+ Dashboards

+ Tags

+ Alerts

+ Services

+ Annotations

+ Chart Tokens

+ Users

Examples & payload structure

CRUD method description

GET /v1/metrics/:name

API VERSION 1.0

Description

Returns information for a specific metric. If time interval search parameters are specified will also include a set of metric measurements for the given time span.

URL

`https://metrics-api.librato.com/v1/metrics/:name`

Method

`GET`

Measurement Search Parameters

If optional `time interval search parameters` are specified, the response includes the set of metric measurements covered by the time interval. Measurements are listed by their originating `source` name if one was specified when the measurement was created. All measurements that were created without an explicit source name are listed with the source name `unassigned`.

`source`

Deprecated: Use `sources` with a single source name, e.g [mysource].

`sources`

If `sources` is specified, the response is limited to measurements from those sources. The `sources` parameter should be specified as an array of source names. The response is limited to the set of sources specified in the array.

Examples

Return the metric named `cpu_temp` with up to four measurements at resolution 60.

```
curl \
-u <user>:<token> \
-X GET \
https://metrics-api.librato.com/v1/metrics/cpu_temp?resolution=60&count=4
```

Response Code

200 OK

Response Headers

** NOT APPLICABLE **

Response Body

```
{
  "type": "gauge",
  "display_name": "cpu_temp",
  "resolution": 60,
  "source": {
    "name": "librato.com": [
      {
        "value": 84.5,
        "time": 1234567890,
        "unit": "Fahrenheit"
      },
      {
        "value": 86.7,
        "time": 1234567950,
        "unit": "Fahrenheit"
      },
      {
        "value": 84.6,
        "time": 1234568010,
        "unit": "Fahrenheit"
      },
      {
        "value": 89.7,
        "time": 1234568070,
        "unit": "Fahrenheit"
      }
    ]
  }
}
```



Short description of the whole domain model

Overview

The central data structure in our engine are **Things**, which are data containers to store all the data generated by and about any physical object. Various **Properties** can be attached to any Thing, and the content of each property can be updated any time, while preserving the history of those changes. Things can be added to various **Collections** which makes it easier to share a set of Things with other **Users** within the engine.

Thing

An abstract notion of an object which has location & property data associated to it. Also called Active Digital Identities (ADIs), these resources can model real-world elements such as persons, places, cars, guitars, mobile phones, etc.

Property

A Thing has various properties: arbitrary key/value pairs to store any data. The values can be updated individually at any time, and can be retrieved historically (e.g. "Give me the values of property X between 10 am and 5 pm on the 16th August 2012").

Location

Each Thing also has a special type of Properties used to store snapshots of its geographic position over time (for now only GPS coordinates - latitude and longitude).

User

Each interaction with the EVRYTHNG back-end is authenticated and a user is associated with each action. This dictates security access.

Creating a new Product

A collection is a grouping of Things. Col one collection.

```
POST /products
Content-Type: application/json
Authorization: $EVRYTHNG_API_KEY

{
  *"fn": <String>,
  "description": <String>,
  "brand": <String>,
  "categories": [<String>, ...],
  "photos": [<String>, ...],
  "url": <String>,
  "identifiers": {
    <String>: <String>,
    ...
  },
  "properties": {
    <String>: <String>,
    ...
  },
  "tags": [<String>, ...]
}
```

Mandatory Parameters

fn

<String> The functional name of the product.

Optional Parameters

description

<String> An string that gives more details about the product, a short description.

CRUD method description

heig-vd

Haute Ecole d'Ingénierie et de Gestion du Canton de Vaud

More details about the Product resource (domain model) & payload structure

Products

Products are very similar to things, but instead of modeling an individual object instance, products are used to model a class of objects. Usually, they are used for general classes of things, usually a particular model with specific characteristics. Let's take for example a specific TV model (e.g. [this one](#)), which has various properties such as a model number, a description, a brand, a category, etc. Products are useful to capture the properties that are common to a set of things (so you don't replicate a property "model name" or "weight" for thousands of things that are individual instances of a same product category).

The Product document model used in our engine has been designed to be compatible with the [hProduct microformat](#), therefore it can easily be integrated with the hProduct data model and applications supporting microformats.

The Product document model is as follows:

```
<Product>={
  "id": <String>,
  "createdAt": <timestamp>,
  "updatedAt": <timestamp>,
  "fn": <String>,
  "description": <String>,
  "brand": <String>,
  "categories": [<String>, ...],
  "photos": [<String>, ...],
  "url": <String>,
  "identifiers": {
    <String>: <String>,
    ...
  },
  "properties": {
    <String>: <String>,
    ...
  },
  "tags": [<String>, ...]
```

Cross-cutting concerns

Pagination

Requests that return multiple items will be paginated to 30 items by default. You can specify further pages with the **?page** parameter. You can also set a custom page size up to 100 with the **?per_page** parameter.

Authentication

Access to our API is done via HTTPS requests to the <https://api.evrythng.com> domain. Unencrypted HTTP requests are accepted (<http://api.evrythng.com> for low-power device without SSL support), but we strongly suggest to use only HTTPS if you store any valuable data in our engine. Every request to our API must include an API key using **Authorization** HTTP header to identify the user or application issuing the request and execute it if authorized.



Instagram

Interactive test console

API Console

Our API console is provided by Apigee. Tap the Lock icon, select OAuth, and you can experiment with making requests to our API. [See it in full screen](#)

The screenshot shows the Instagram API console interface. At the top, there's a header with the Instagram logo and the text "Interactive test console". Below that is a section titled "API Console" with a sub-section "Service" containing the URL "https://api.instagram.com/v1" and an "Authentication" dropdown set to "No Auth". The main area is titled "Select an API method" and contains two sections: "Users" and "Relationships". Under "Users", there are five methods: GET users/(user-id), GET users/self/feed, GET users/(user-id)/media/recent, GET users/self/media/liked, and GET users/search. Under "Relationships", there are four methods: GET users/(user-id)/follows, GET users/(user-id)/followed-by, GET users/self/requested-by, and GET users/(user-id)/relationship (GET). Each method has a blue button next to it.

GET /media/ media-id /comments

https://api.instagram.com/v1/media/555/comments?access_token=ACCESS-TOKEN

RESPONSE

```
{
  "meta": {
    "code": 200
  },
  "data": [
    {
      "created_time": "1280780324",
      "text": "Really amazing photo!",
      "from": {
        "username": "snooppogg",
        "profile_picture": "http://images.instagram.com/profiles/profile_16_75sq_1305612434.jpg",
        "id": "1574083",
        "full_name": "Snoop Dogg"
      },
      "id": "420"
    },
    ...
  ]
}
```

Get a full list of comments on a media.
Required scope: comments

CRUD method description

List of supported CRUD methods for each resource (R, RAW)

Haute Ecole d'Ingénierie et de Gestion du canton de Vaud

User Endpoints

GET /users/ user-id	... Get basic information about a user.
GET /users/self/feed	... See the authenticated user's feed.
GET /users/ user-id /media/recent	... Get the most recent media published by a user.
GET /users/self/media/liked	... See the authenticated user's list of liked media.
GET /users/search	... Search for a user by name.

Comment Endpoints

GET /media/ media-id /comments	... Get a full list of comments on a media.
POST /media/ media-id /comments	... Create a comment on a media. Please email apide...
DEL /media/ media-id /comments/ comment-id	... Remove a comment.

Cross-cutting concerns

Limits

Be nice. If you're sending too many requests too quickly, we'll send back a 503 error code (server unavailable).

You are limited to 5000 requests per hour per access_token or client_id overall. Practically, this means you should (when possible) authenticate users so that limits are well outside the reach of a given user.

PAGINATION

Sometimes you just can't get enough. For this reason, we've provided a convenient way to access more data in any request for sequential data. Simply call the url in the next_url parameter and we'll respond with the next set of data.

The Envelope

Every response is contained by an envelope. That is, each response has a predictable set of keys with which you can expect to interact:

```
{
  "meta": {
    "code": 200
  },
  "data": [
    ...
  ],
  "pagination": {
    "next_url": "...",
    "next_max_id": "13872296"
  }
}
```



Tools & Demo